

## Testing Conditions

- **CPU**
- Results are from `run_rand_bench`. Each trials (**100,000**) runs **4096** iterations the `al_src`, `al_dst`, and `sz` parameters. These values are taken from the SPEC2017 distribution. For purpose of simplifying these results `sz` is only of values in range `[0, 32]`.
- The memcpy implementation I tested only supported `sz` in `[0, 32]` in an attempt to isolate what I believe to be the problem.
- The results are from **throughput** tests. **NOT latency**.
- All of the benchmark code is on its own page and only uses address range `% 4096` in `[0, 2047]`
- The memcpy code is on its own page and only uses address range `% 4096` in `[2048, 4095]`.

## Notes on Testing Conditions

- **The testing conditions described above are the only ones I have run the entire benchmark for** (roughly 16 hours). I have tinkered with some other conditions that I though may have had an effect or be skewing the results entirely and **Did NOT see any conditions which appeared to alter the trends described below**.
- The benchmark loop calling memcpy is aligned to 64 bytes. I have not run the full benchmark on other loop alignments but have somewhat tested that other alignments / addition `nop` padding values do not change the trend in the results (although can have some effect on the exact values by what seems to be a constant factor).
- Splitting the address space on different boundaries, for example benchmark code in `[0, 511]` and memcpy in `[512, 4095]` (or `[0, 1023]` and `[1024, 4095]`) also yielded the same trend with minimal (if any) effect on the results.
- Measuring latency as opposed to throughput significantly shrunk the variance between the “fast” and “slow” mode, although the same trend in overall performance and performance counter values was still visible.

## Notes on Tested Parameters / Implementation

The **parameters** tested are as follows: - `ALIGN_ENTRY` in `[16, 32, 48, 64]`.  
- The formula for computing the entry address `% 4096` of memcpy is `2048 + (ALIGN_ENTRY % 64)`. - `PADDING` in `[0...1025]` s.t padding `% 16` in `[0, 1]`

The **implementation** with the above parameters is as follows:

```
// + ALIGN_ENTRY % 64
```

```
memcpy:
```

```
800:  48 83 fa 20          cmp     $0x20,%rdx
804:  72 2a              jb     830 <memcpy_dev_v32_movsb+0x30>
```

```

806: 62 e1 fe 28 6f 06      vmovdqu64 (%rsi),%ymm16
80c: 62 61 fe 28 6f 7c 16   vmovdqu64 -0x20(%rsi,%rdx,1),%ymm31
813: ff
814: 62 e1 fe 28 7f 07      vmovdqu64 %ymm16,(%rdi)
81a: 62 61 fe 28 7f 7c 17   vmovdqu64 %ymm31,-0x20(%rdi,%rdx,1)
821: ff
822: c3                     retq
823: 66 66 2e 0f 1f 84 00   data16 nopw %cs:0x0(%rax,%rax,1)
82a: 00 00 00 00
82e: 66 90                  xchg    %ax,%ax

// + PADDING

L(less_vec):
830: 83 fa 10               cmp     $0x10,%edx
833: 73 2b                  jae     860 <memcpy_dev_v32_movsb+0x60>
835: 83 fa 08               cmp     $0x8,%edx
838: 73 46                  jae     880 <memcpy_dev_v32_movsb+0x80>
83a: 83 fa 04               cmp     $0x4,%edx
83d: 73 61                  jae     8a0 <memcpy_dev_v32_movsb+0xa0>
83f: 83 fa 01               cmp     $0x1,%edx
842: 72 11                  jb      855 <memcpy_dev_v32_movsb+0x55>
844: 0f b6 4c 16 ff         movzbl -0x1(%rsi,%rdx,1),%ecx
849: 74 06                  je      851 <memcpy_dev_v32_movsb+0x51>
84b: 0f b7 36               movzwl (%rsi),%esi
84e: 66 89 37               mov     %si,(%rdi)
851: 88 4c 17 ff            mov     %cl,-0x1(%rdi,%rdx,1)
855: c3                     retq
856: 66 2e 0f 1f 84 00 00   nopw   %cs:0x0(%rax,%rax,1)
85d: 00 00 00
860: c5 fa 6f 06           vmovdqu (%rsi),%xmm0
864: c5 fa 6f 4c 16 f0     vmovdqu -0x10(%rsi,%rdx,1),%xmm1
86a: c5 fa 7f 07           vmovdqu %xmm0,(%rdi)
86e: c5 fa 7f 4c 17 f0     vmovdqu %xmm1,-0x10(%rdi,%rdx,1)
874: c3                     retq
875: 66 66 2e 0f 1f 84 00   data16 nopw %cs:0x0(%rax,%rax,1)
87c: 00 00 00 00
880: 48 8b 0e              mov     (%rsi),%rcx
883: 48 8b 74 16 f8        mov     -0x8(%rsi,%rdx,1),%rsi
888: 48 89 0f              mov     %rcx,(%rdi)
88b: 48 89 74 17 f8        mov     %rsi,-0x8(%rdi,%rdx,1)
890: c3                     retq
891: 66 66 2e 0f 1f 84 00   data16 nopw %cs:0x0(%rax,%rax,1)
898: 00 00 00 00
89c: 0f 1f 40 00           nopl    0x0(%rax)
8a0: 8b 0e                 mov     (%rsi),%ecx

```

```

8a2:  8b 74 16 fc          mov     -0x4(%rsi,%rdx,1),%esi
8a6:  89 0f                mov     %ecx,(%rdi)
8a8:  89 74 17 fc          mov     %esi,-0x4(%rdi,%rdx,1)
8ac:  c3                  retq


```



### Please Note


- The exact address, and address % 64 of L(less\_vec) changes with ALIGN\_ENTRY.
- PADDING essentially represents address\_of(memcpy) - address\_of(L(less\_vec)) + 48

### Results

The results for Ref Cycles (from rdtsc) vs PADDING for ALIGN\_ENTRY in [16, 64] are:







The results seems to depend on `ALIGN_ENTRY` and `PADDING % 64`. The following table summarizes the basic results:

ALIGN_ENTRY	0	16	32	48
16	SLOW	SLOW	FAST	FAST
32	FAST	SLOW	SLOW	FAST
48	SLOW	SLOW	FAST	FAST
64	FAST	SLOW	SLOW	FAST